

# Log Monitoring Architecture

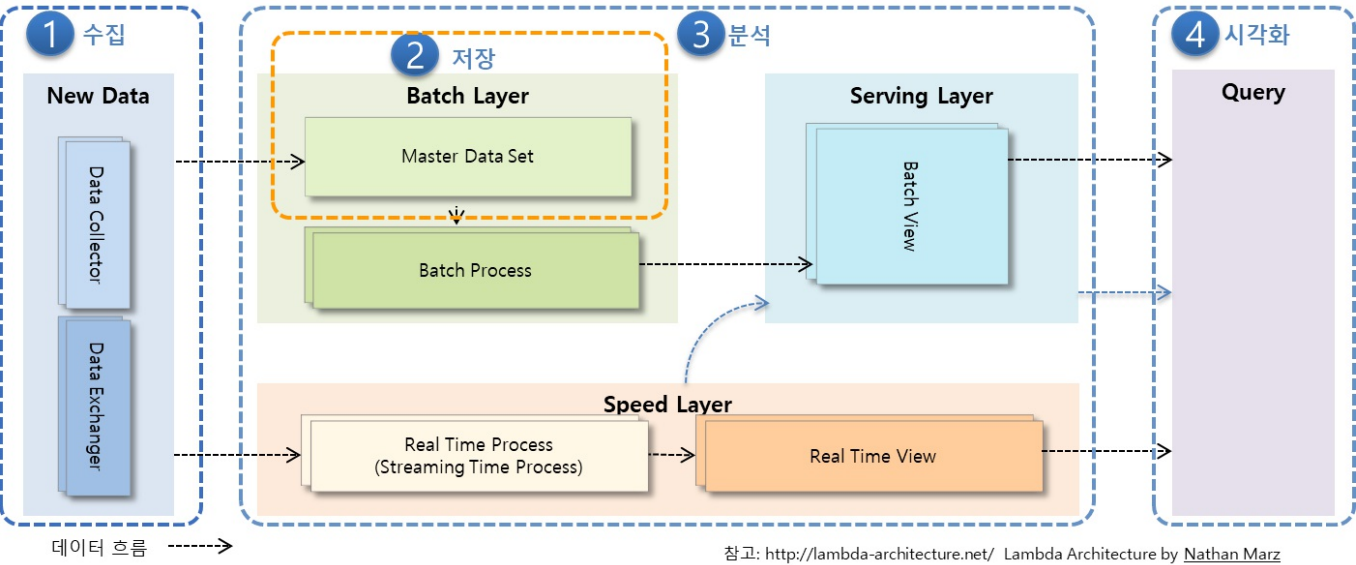
## 목차

- 1. 모니터링 아키텍처 및 요소기술
- 2. 하둡 및 HBase를 이용한 대용량 테이블 설계
- 3. Storm 을 이용한 실시간 데이터 분석
- 4. Kafka 및 Redis 활용
- 5. Grafana 와 InfluxDB를 이용한 대시보드 구성

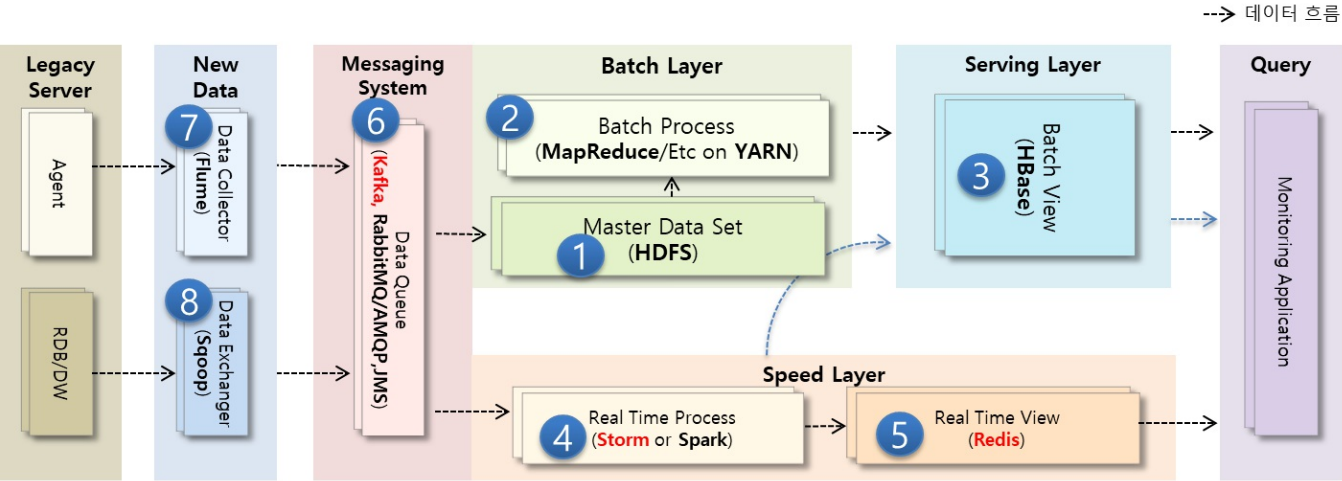
## 모니터링 아키텍처 및 요소기술

### Lambda Architecture

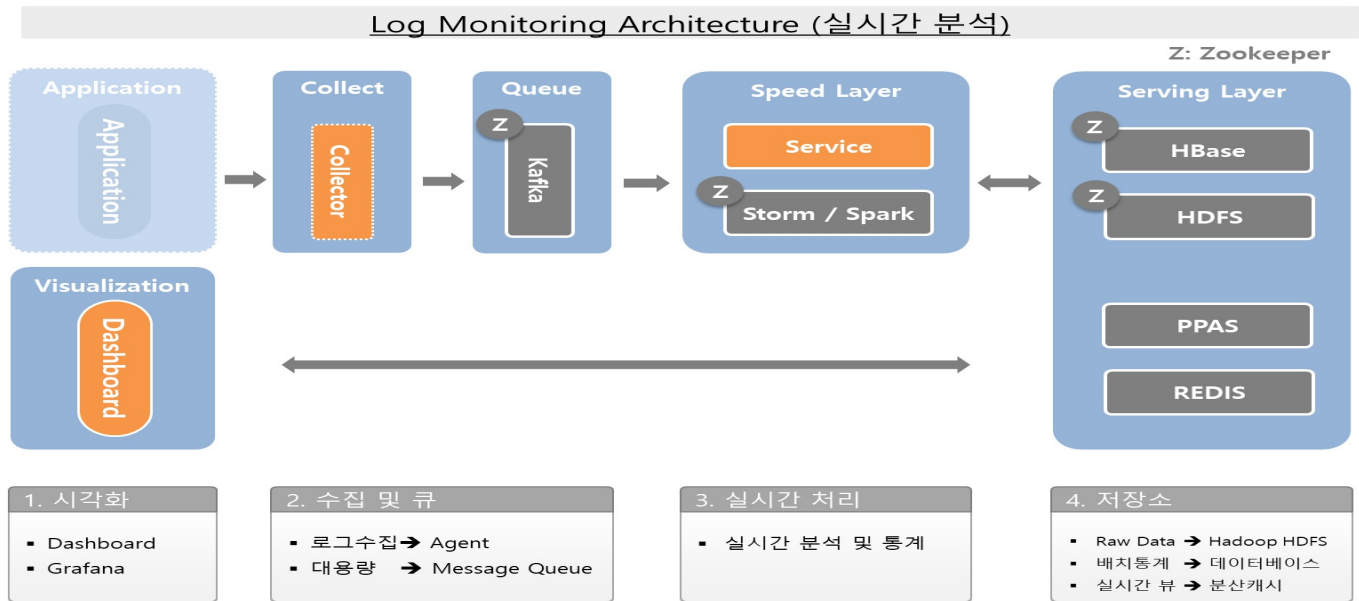
실시간 분석을 지원하는 빅데이터 아키텍처



람다 아키텍처에 대한 솔루션 매핑 사례



## Monitoring Architecture(Reference)



## Tech Stack

1. Distributed Platform: Hadoop
2. Data Store: **HBase**
3. Distribute Codinator: Zookeeper
4. Realtime Analysis: **Storm** or Spark
5. Visualization: Grafana
6. TSDB: InfluxDB
7. Collect: Telegraf
8. Message Queue: **Kafka**
9. In Memory Data Grid: Redis

Other Hadoop-related projects at Apache include:

- Ambari
  - A web-based tool for provisioning, managing, and monitoring Apache Hadoop clusters which includes support for Hadoop HDFS, Hadoop MapReduce, Hive, HCatalog, HBase, ZooKeeper, Oozie, Pig and Sqoop. Ambari also provides a dashboard for viewing cluster health such as heatmaps and ability to view MapReduce, Pig and Hive applications visually alongwith features to diagnose their performance characteristics in a user-friendly manner.
- Avro
  - A data serialization system
  - cf. **Thrift**, Protocol Buffers
- Hive
  - A data warehouse infrastructure that provides data summarization and ad hoc querying.
- Pig
  - A high-level data-flow language and execution framework for parallel computation.
- Spark
  - A fast and general compute engine for Hadoop data. Spark provides a simple and expressive

programming model that supports a wide range of applications, including ETL, machine learning, stream processing, and graph computation.

## 하둡 및 HBase를 이용한 대용량 테이블 설계

### 2.1 Hadoop

The Apache Hadoop software library is a framework that allows for the distributed processing of large data sets across clusters of computers using simple programming models.

- reliable, scalable, distributed computing

#### HDFS(Hadoop Distributed File System)

The Hadoop Distributed File System (HDFS) is a distributed file system designed to run on commodity hardware.

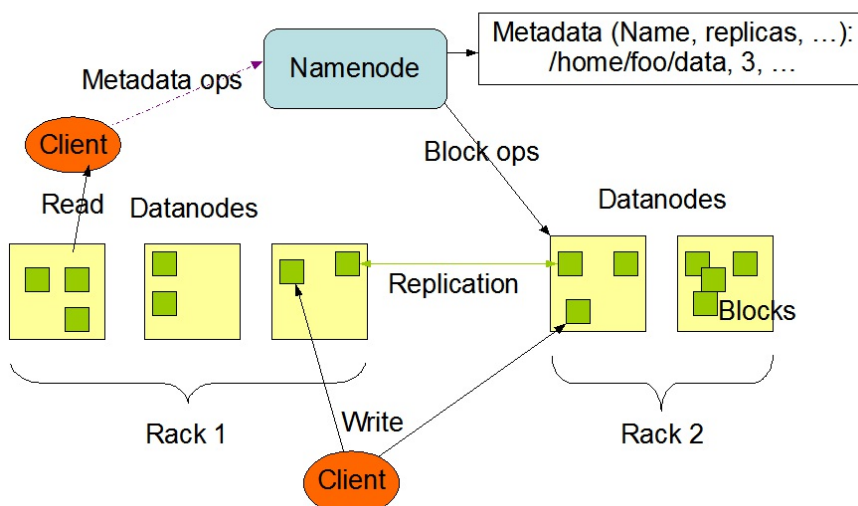
#### Assumptions and Goals

- Hardware Failure
- Streaming Data Access
  - Write Once, Read Any Number Of Times, But Don't Try To Change The Contents Of The File.
  - Throughput > Latency
- Large Data Sets
- Simple Coherency Model
- "Moving Computation is Cheaper than Moving Data"
- Portability Across Heterogeneous Hardware and Software Platforms

#### HDFS Architecture

Master-Slave architecture

HDFS Architecture



**NameNode**

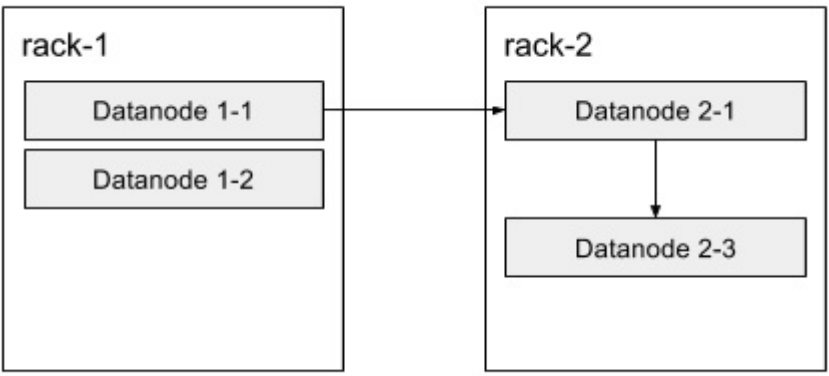
- 파일 시스템의 네임스페이스를 관리
- 클라이언트의 파일 접근을 관리
- 파일 이름 변경, 파일과 디렉토리의 오픈/클로즈와 같은 파일 시스템 작업을 실행

**DataNode**

- 클라이언트의 요청을 읽어서 파일에 읽기/쓰기 작업
- 블록의 생성, 삭제, 복제와 같은 작업을 수행

**Block**

- 여러 데이터 노드에 작은 파일로 나뉘다음 분산되어 저장
- 블록의 기본크기는 64MB, HDFS 설정에 따라서 바꿀 수 있음
- 블록을 보관할 노드의 선택
  - 첫번째 복제는 원본과 같은 랙에 있는 노드를 선택
  - 두번째와 세번째는 다른 랙에 보관



**Operations**

Action	Command
Create a directory named /foodir	bin/hadoop dfs -mkdir /foodir
Remove a directory named /foodir	bin/hadoop fs -rm -R /foodir
View the contents of a file named /foodir/myfile.txt	bin/hadoop dfs -cat /foodir/myfile.txt

Hadoop 기본 포트

파라미터	기본 값	용도	설정 파일
dfs.secondary.http.address	50090	SecondaryNameNode 웹 서버 주소 및 포트	hdfs-site.xml
dfs.datanode.address	50010		

DataNode 주소 및 포트			hdfs-site.xml
dfs.datanode.http.address	50075	DataNode 웹 서버 용 주소 및 포트	hdfs-site.xml
dfs.datanode.ipc.address	50020	DataNode ipc 서버용 주소 및 포트	hdfs-site.xml
dfs.http.address	50070	NameNode 웹 어드민용 주소 및 포트	hdfs-site.xml
mapred.job.tracker.http.address	50030	JobTracker 웹 서버용 주소 및 포트	mapred-site.xml
mapred.task.tracker.http.address	50060	TaskTracker 웹 서버용 웹 주소 및 포트	mapred-site.xml

## Apache Hadoop API

### pom.xml

```
<dependency>
  <groupId>org.apache.hadoop</groupId>
  <artifactId>hadoop-client</artifactId>
  <version>2.7.1</version>
</dependency>
```

### HDFSClient.java

```
System.setProperty("HADOOP_USER_NAME", "vagrant");
UserGroupInformation ugi = UserGroupInformation.createRemoteUser("vagrant");

Configuration conf = new Configuration();
conf.set("fs.default.name", "hdfs://hadoop01:9000");

FileSystem dfs = FileSystem.get(conf);

System.out.println("Home Path : " + dfs.getHomeDirectory());
System.out.println("Work Path : " + dfs.getWorkingDirectory());

Path filenamePath = new Path("/tmp/hello.txt");
System.out.println("File Exists : " + dfs.exists(filenamePath));

if (dfs.exists(filenamePath)) {
    dfs.delete(filenamePath, true);
}

FSDataOutputStream out = dfs.create(filenamePath);
out.writeUTF("Hello, world!\n");
out.close();

FSDataInputStream in = dfs.open(filenamePath);
```

```
String messageIn = in.readUTF();
System.out.print(messageIn);

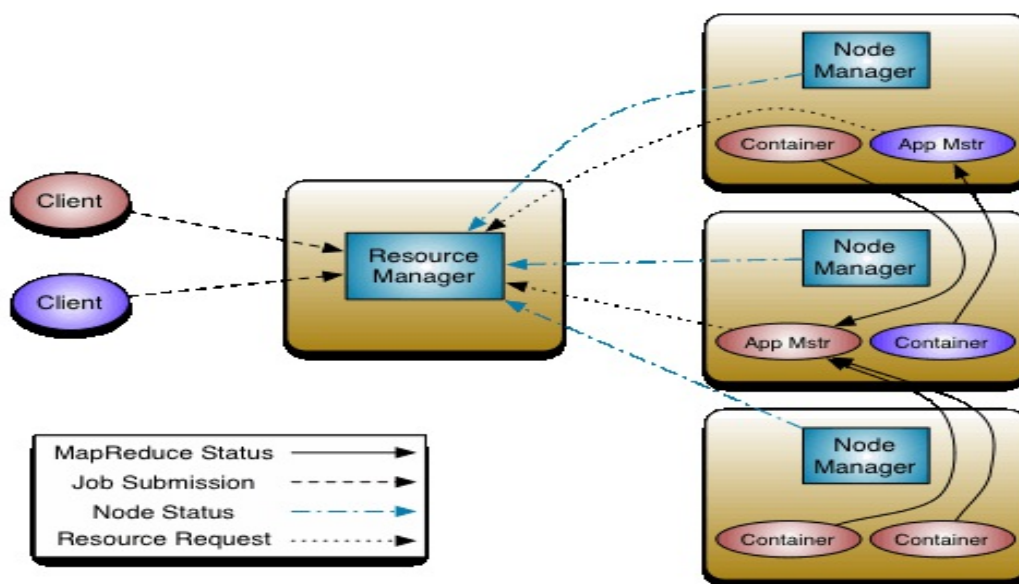
in.close();
dfs.close();
```

## 실행결과

```
[vagrant@hadoop03 ~]$ hadoop fs -ls /tmp
Found 1 items
-rw-r--r--  3 vagrant supergroup      16 2018-01-24 19:47 /tmp/hello.txt
[vagrant@hadoop03 ~]$ hadoop fs -cat /tmp/hello.txt
Hello, world!
```

## Map & Reduce

### Apache Hadoop NextGen MapReduce (YARN)



## Functional Programming

### Map

- Data Processing

### Reduce

- Merge

## 2.2 HBase

Apache HBase™ is the Hadoop database, a distributed, scalable, big data store.

- NoSQL
- Schemaless
- 가용성과 확장성
  - Region 서버만 추가하여 확장성 및 가용성 확보가 용이
  - 단, 특정 Region 서버에 부하가 집중되면 성능저하가 발생 -> 로우키 설계가 중요
- CAP: "분산 시스템에서는 다음의 3개 속성을 모두 가지는 것이 불가능하다"
  - Consistency (V)
  - **Availability**
  - Partitions Tolerance (V)
- Apache Phoenix: Add SQL Layer

## Column Family

The diagram illustrates a table structure with two column families: 'Customer' and 'Sales'. The 'Customer' family contains 'Name' and 'City' columns, while the 'Sales' family contains 'Product' and 'Amount' columns. The 'Row Key' column is used to identify each row. Annotations include: 'cells' pointing to a specific data cell (Jonh White), 'Column Family' pointing to the Sales family, 'Column' pointing to the Amount column, and 'ROW' pointing to the row with Row Key 102.

Row Key	Customer		Sales	
customer id	Name	City	Product	Amount
101	Jonh White	Los Angeles	Chairs	400.00
102	Jane Brown	Atlanta	Lamps	200.00
103	Bill Green	Pittsburgh	Desk	500.00
104	Jack Black	St. Louis	Bed	1600.00

1. 이 테이블은 Customer와 Sales 두 개의 컬럼 패밀리를 가지고 있다.
2. Customer 컬럼 패밀리는 Name과 City 두 개의 컬럼을 가지고 있다.
3. Sales 컬럼 패밀리는 Product와 Amount 두 개의 컬럼을 가지고 있다.
4. Row는 Row Key, Customer CF, Sales CF로 구성된다.

## 특징

row-key => {

cf1:column-key1: 0 (at timestamp t1),  
cf1:column-key1: 1 (at timestamp t2),

"Expired" data with older timestamps, not returned by default

cf1:column-key1: 2 (at timestamp t3),

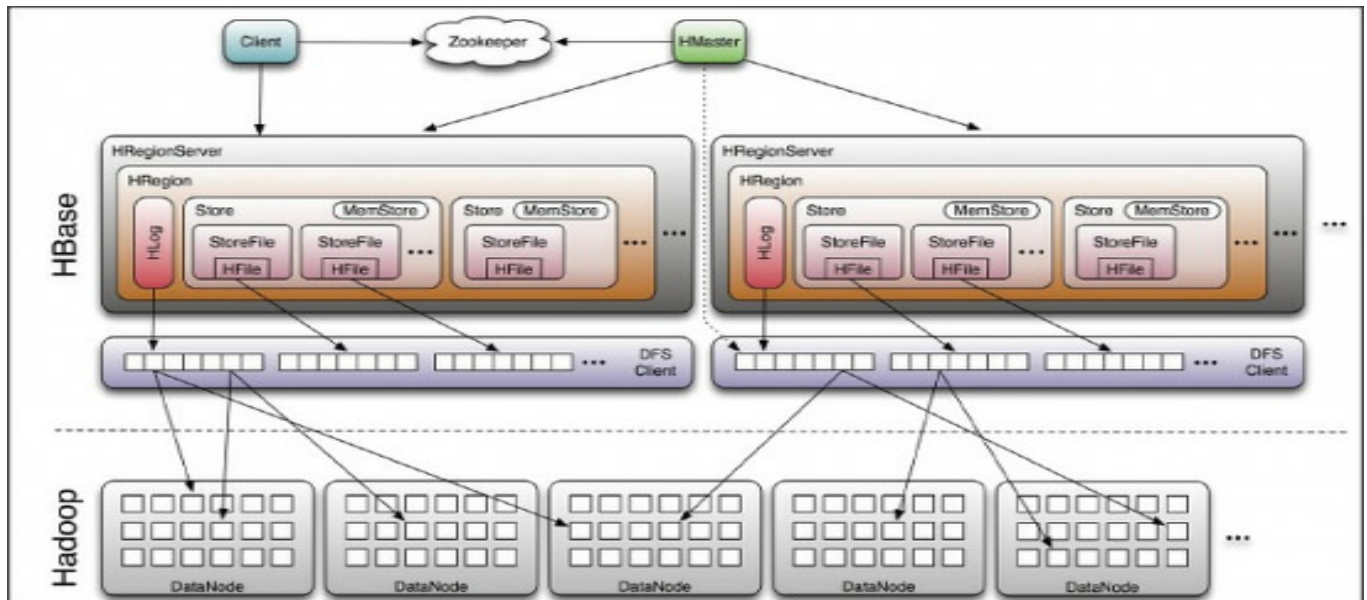
"Current" or "live" value which is returned on a query

.....  
}

Row Key	Timestamp	Customer		Sales	
		Name	City	Product	Amount
101	T1	Suresh	Hyderabad		300.00
101	T2	Suresh Reddy		Books	
102	T1	Lavya Gavshinde	Indore	Fan	600
102	T2		Bhopal		
102	T3	Deepesh	Delhi	Bike	32000

1. Sparse: HBase는 sparse matrix(희소행렬)방식으로 데이터를 저장 할 수 있다. 예컨데 굳이 모든 필드에 값을 채울 필요가 없다는 얘기가.
2. Column Oriented: RDBMS는 row 단위로 데이터를 저장한다. 하지만 HBase는 Sparse하기 때문에, 컬럼 단위로 데이터를 읽고 쓸 수 있다.
3. Distributed: 테이블은 수백만개의 Row와 컬럼들로 구성된다. 이들 컬럼들은 쉽게 분산 할 수 있다.
4. Versioned: 위 테이블은 Customer ID를 Row key로 사용하고 있다. 당연하지만 고객은 하나 이상의 물건을 구매할 수 있기 때문에 하나의 row key에 여러 개의 row를 저장 할 수 있다. 이 경우 Timestamp가 버전이 되며, 질의를 할 경우 가장 최근 timestamp의 row를 읽어 온다.
5. Non-relational: 관계형 데이터베이스의 ACID 속성을 따르지 않는다.

## Architecture



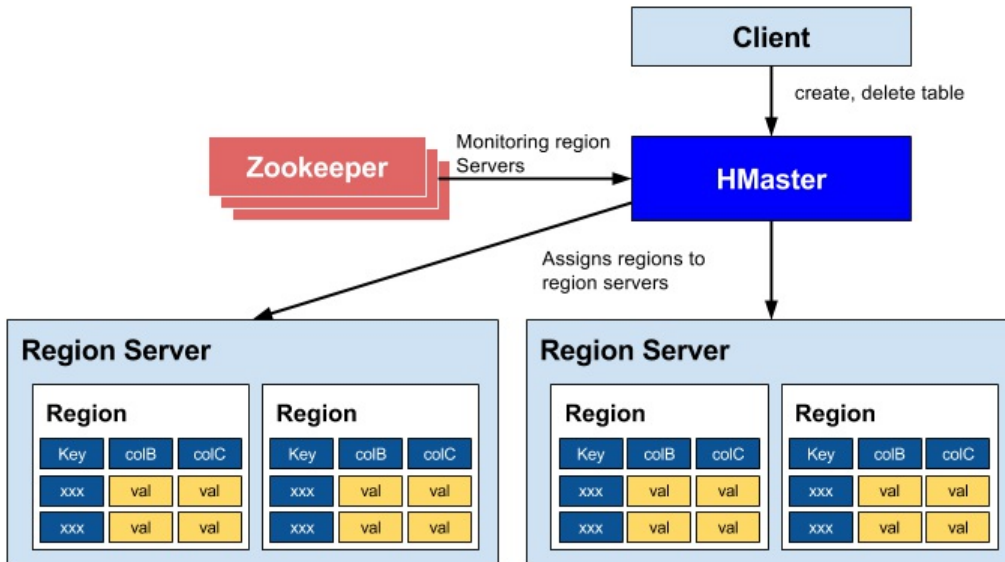
- HMaster
  - RegionServer 와 META 데이터 관리
- HRegionServer
  - Region 관리



- HResion
  - 테이블 데이터의 부분집합이며, 수평확장의 기본단위
- HLog
- HFile

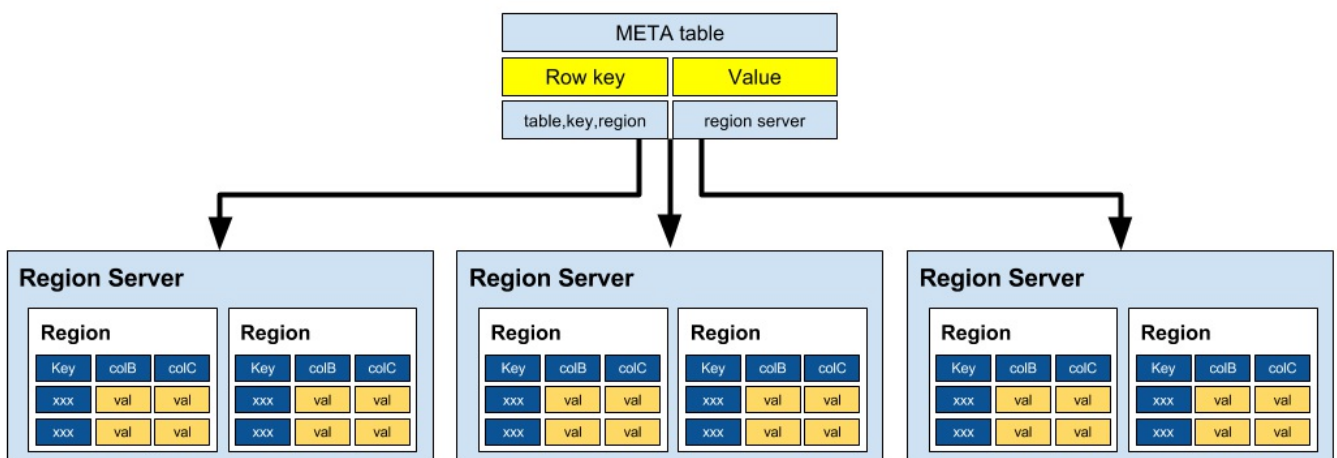
## HBase HMaster

각 테이블의 데이터는 HRegionServer가 관리하며, 전체 클러스터는 HMaster가 관리한다.



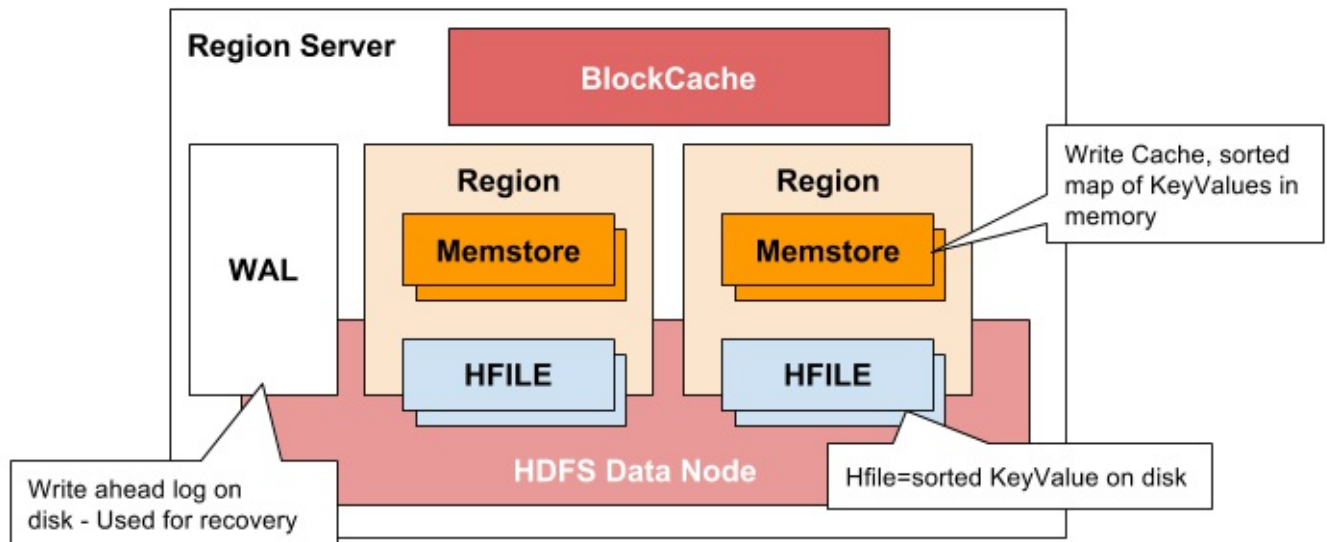
- 리전 서버들을 조정
  - 리전의 시작을 관리
  - 클러스터에 있는 모든 리전 서버들을 모니터링
- 관리 기능
  - 테이블의 생성, 삭제, 업데이트

## HBase Catalog Tables



- META 테이블
  - 클러스터에 포함된 리전의 위치정보들을 저장

## Region Server Components



- WAL(Write Ahead Log)
  - 데이터 저장 실패를 복구하기 위해서 사용
- BlockCache
  - 읽기 캐시
- MemStore
  - 쓰기 캐시
  - 각 리전의 컬럼 패밀리당 하나
- HFile

## HBASE vs RDBMS

HBase	RDBMS
스키마가 없다. column families만으로 이용	테이블의 구조를 기술하는 스키마
수평적으로 확장성이 있어 큰 테이블에 적합	확장하기 어려우며, 크기가 작은 테이블을 위해 생성
Transaction이 존재하지 않음	Transaction이 존재
비 일반화된 데이터	일반화된 데이터
덜 구조화된 데이터가 적합	구조화된 데이터에 적합
get / put / scan 등	SQL
MapReduce Join 활용	Join에 최적화 됨
RowKey만 인덱스 지원	임의 컬럼에 대한 인덱스 지원
초당 수십만건 Read/Write	초당 수천건 Read/Write
단일로우 트랜잭션 보장	다중 로우 트랜잭션 보장

## HBase 사용법

## 1. Connect to HBase

```
$ ./bin/hbase shell
hbase(main):001:0>
```

## 2. Create a table

```
hbase(main):001:0> create 'test', 'cf'
0 row(s) in 0.4170 seconds

=> Hbase::Table - test
```

## 3. List Information About your Table

```
hbase(main):002:0> list 'test'
TABLE
test
1 row(s) in 0.0180 seconds

=> ["test"]
```

## 4. Put data into your table.

```
hbase(main):003:0> put 'test', 'row1', 'cf:a', 'value1'
0 row(s) in 0.0850 seconds

hbase(main):004:0> put 'test', 'row2', 'cf:b', 'value2'
0 row(s) in 0.0110 seconds

hbase(main):005:0> put 'test', 'row3', 'cf:c', 'value3'
0 row(s) in 0.0100 seconds
```

## 5. Scan the table for all data at once.

```
hbase(main):006:0> scan 'test'
ROW                                COLUMN+CELL
  row1                             column=cf:a, timestamp=1421762485768,
value=value1
  row2                             column=cf:b, timestamp=1421762491785,
value=value2
```

```
row3                                     column=cf:c, timestamp=1421762496210,  
value=value3  
3 row(s) in 0.0230 seconds
```

## 6. Get a single row of data.

```
hbase(main):007:0> get 'test', 'row1'  
COLUMN                                CELL  
cf:a                                  timestamp=1421762485768, value=value1  
1 row(s) in 0.0350 seconds
```

## 7. Disable a table.

```
hbase(main):008:0> disable 'test'  
0 row(s) in 1.1820 seconds  
  
hbase(main):009:0> enable 'test'  
0 row(s) in 0.1770 seconds
```

## 8. Drop the table.

```
hbase(main):011:0> drop 'test'  
0 row(s) in 0.1370 seconds
```

## 9. Exit the HBase Shell.

```
hbase(main):006:0> quit
```

## HBase Java Client Example

### HBase Data Structure

```
Family1:{  
  'Qualifier1':'row1:cell_data',  
  'Qualifier2':'row2:cell_data',  
  'Qualifier3':'row3:cell_data'  
}  
Family2:{  
  'Qualifier1':'row1:cell_data',
```

```
'Qualifier2':'row2:cell_data',  
'Qualifier3':'row3:cell_data'  
}
```

## HBase Client Maven Dependency

```
<dependency>  
  <groupId>org.apache.hbase</groupId>  
  <artifactId>hbase-client</artifactId>  
  <version>${hbase.version}</version>  
</dependency>  
<dependency>  
  <groupId>org.apache.hbase</groupId>  
  <artifactId>hbase</artifactId>  
  <version>${hbase.version}</version>  
  <type>pom</type>  
</dependency>
```

## 1. Connecting to HBase

hbase-site.xml

```
<configuration>  
  <property>  
    <name>hbase.zookeeper.quorum</name>  
    <value>hadoop01</value>  
  </property>  
  <property>  
    <name>hbase.zookeeper.property.clientPort</name>  
    <value>2181</value>  
  </property>  
</configuration>
```

```
Configuration config = HBaseConfiguration.create();  
  
String path = this.getClass()  
    .getClassLoader()  
    .getResource("hbase-site.xml")  
    .getPath();  
config.addResource(new Path(path));  
  
HBaseAdmin.checkHBaseAvailable(config);
```

## 2. Creating a Database Structure

```
private TableName table1 = TableName.valueOf("Table1");
private String family1 = "Family1";
private String family2 = "Family2";
```

```
Connection connection = ConnectionFactory.createConnection(config)
Admin admin = connection.getAdmin();
```

```
HTableDescriptor desc = new HTableDescriptor(table1);
desc.addFamily(new HColumnDescriptor(family1));
desc.addFamily(new HColumnDescriptor(family2));
admin.createTable(desc);
```

### 3. Adding and Retrieving Elements

```
byte[] row1 = Bytes.toBytes("row1")
Put p = new Put(row1);
p.addImmutable(family1.getBytes(), qualifier1, Bytes.toBytes("cell_data"));
table1.put(p);
```

```
Get g = new Get(row1);
Result r = table1.get(g);
byte[] value = r.getValue(family1.getBytes(), qualifier1);
```

```
Bytes.toString(value)
```

### 4. Scanning and Filtering

```
Filter filter1 = new PrefixFilter(row1);
Filter filter2 = new QualifierFilter(
    CompareOp.GREATER_OR_EQUAL,
    new BinaryComparator(qualifier1));
List<Filter> filters = Arrays.asList(filter1, filter2);
```

```
Scan scan = new Scan();
```

```
scan.setFilter(new FilterList(Operator.MUST_PASS_ALL, filters));
scan.addColumn(family1.getBytes(), qualifier1);

ResultScanner scanner = table.getScanner(scan);
for (Result result : scanner) {
    System.out.println("Found row: " + result);
}
```

## 5. Deleting Rows

```
Delete delete = new Delete(row1);
delete.addColumn(family1.getBytes(), qualifier1);
table.delete(delete);
```

---

## Apache Phoenix

Phoenix is an open source SQL skin for HBase.


You use the standard JDBC APIs instead of the regular HBase client APIs to create tables, insert data, and query your HBase data.

### 특징

- HBase 데이터에 대한 빠른 접근이 가능
  - MapReduce를 사용하지 않고 SQL 쿼리를 Native HBase 명령어로 컴파일하여 수행
- 추가적인 서버가 필요하지 않기 때문에 비교적 가벼움
- 기존에 존재하는 HBase 테이블과의 매핑이 가능

### 장점

- JOIN
- Paging
  - OFFSET & LIMIT (ex. LIMIT 10 OFFSET 10)
- Secondary Index
- Salted Tables
  - Splitting: pre-split on salt bytes boundaries
    - ex. new\_row\_key = (++index % BUCKETS\_NUMBER) + original\_key



- Row Key Ordering: phoenix.query.rowKeyOrderSaltedTable=true
- CREATE TABLE table (a\_key VARCHAR PRIMARY KEY, a\_col VARCHAR) SALT\_BUCKETS = 20;

## Apache Phoenix Java Client Example

### Phoenix Driver Maven Dependency

```
<dependency>
  <groupId>org.apache.phoenix</groupId>
  <artifactId>phoenix-core</artifactId>
  <version>4.13.1-HBase-1.2</version>
</dependency>
```

### PhoenixExample.java

```
public class PhoenixExample {

    public static void main(String[] args) {
        // Create variables
        Connection connection = null;
        Statement statement = null;
        ResultSet rs = null;
        PreparedStatement ps = null;

        try {
            connection =
DriverManager.getConnection("jdbc:phoenix:demo");
            statement = connection.createStatement();
            statement.executeUpdate("create table javatest (mykey
integer not null primary key, mycolumn varchar)");
            statement.executeUpdate("upsert into javatest values
(1,'Hello')");
            statement.executeUpdate("upsert into javatest values
(2,'Java Application')");
            connection.commit();

            ps = connection.prepareStatement("select * from
javatest");
```



```

        rs = ps.executeQuery();

        System.out.println("Table Values");
        while (rs.next()) {
            Integer myKey = rs.getInt(1);
            String myColumn = rs.getString(2);
            System.out.println("\tRow: " + myKey + " = " +
myColumn);
        }
    } catch (SQLException e) {
        e.printStackTrace();
    } finally {
        try {
            ps.close();
            rs.close();
            statement.close();
            connection.close();
        } catch (Exception e) {
        }
    }
}
}
}


```

## 2.3 ZooKeeper

A Distributed Coordination Service for Distributed Applications

- 설정 정보 관리
- 노드 리스트 관리
- 리더선정

### ZooKeeper Architecture




### Design Goals

- ZooKeeper is simple.
  - ZooKeeper allows distributed processes to coordinate with each other through a shared hierarchal namespace which is organized similarly to a standard file system.

- Zookeeper is replicated.
  - ZooKeeper itself is intended to be replicated over a sets of hosts called an ensemble.
- Zookeeper is ordered.
  - ZooKeeper stamps each update with a number that reflects the order of all ZooKeeper transactions.
- Zookeeper is fast.
  - It is especially fast in "read-dominant" workloads.

## Data model and the hierarchical namespace



## Nodes and ephemeral nodes

- each node in a ZooKeeper namespace can have data associated with it as well as children.
- These znodes exists as long as the session that created the znode is active.

## Simple API

operation	desc
create	creates a node at a location in the tree
delete	deletes a node
exists	tests if a node exists at a location
get data	reads the data from a node
set data	writes data to a node
get children	retrieves a list of children of a node
sync	waits for data to be propagated

## Example

### Maven Dependency

```
<dependency>
  <groupId>org.apache.zookeeper</groupId>
  <artifactId>zookeeper</artifactId>
  <version>3.4.8</version>
</dependency>
```

## ZNodeTest.java

```
public class ZNodeTest {

    public static void main(String[] args) {
        int sessionTimeout = 10 * 100;

        // 1. ZooKeeper 서버(클러스터)에 연결
        ZooKeeper zk = null;

        try {
            zk = new ZooKeeper("demo:2181", sessionTimeout, null);

            // 2. Test01 znode가 존재하지 않으면 /test01, /test02 생성
            if (zk.exists("/test01", false) == null) {
                zk.create("/test01", "test01_data".getBytes(),
                    Ids.OPEN_ACL_UNSAFE, CreateMode.PERSISTENT);
                zk.create("/test02", "test02_data".getBytes(),
                    Ids.OPEN_ACL_UNSAFE, CreateMode.PERSISTENT);
            }

            // 3 /test01 노드의 자식 노드로 sub01, sub02 생성
            zk.create("/test01/sub01", null, Ids.OPEN_ACL_UNSAFE,
                CreateMode.EPHEMERAL);
            zk.create("/test01/sub02", null, Ids.OPEN_ACL_UNSAFE,
                CreateMode.EPHEMERAL);

            // 4. /test01 노드의 데이터 가져오기
            byte[] test01Data = zk.getData("/test01", false, null);
            System.out.println("getData [/test01]: " + new
                String(test01Data));

            // 5. /test01/sub01 노드의 데이터를 새로운 값으로 설정
            zk.setData("/test01/sub01", "this new Data".getBytes(), -1);
            byte[] subData = zk.getData("/test01/sub01", false, null);
            System.out.println("getData after setData [/test01/sub01]: " + new
                String(subData));

            // 6. 노드가 존재하는지 확인
            System.out.println("exist [/test01/sub01]: " +
                (zk.exists("/test01/sub01", false) != null));
            System.out.println("exist [/test01/sub03]: " +
                (zk.exists("/test01/sub03", false) != null));
```

```

        // 7. /test01의 자식노드 목록 가져오기
        List<String> children = zk.getChildren("/test01", false);
        for (String eachChildren : children) {
            System.out.println("getChildren [/test01]: " + eachChildren);
        }
    } catch (IOException | KeeperException | InterruptedException e) {
        e.printStackTrace();
    } finally {
        // 8. zk 종료
        try {
            zk.close();
        } catch (InterruptedException e) {
            e.printStackTrace();
        }
    }
}
}
}

```

## Storm 을 이용한 실시간 데이터 분석

Apache Storm is a free and open source distributed realtime computation system.

### 데이터 스트림과 실시간 분석

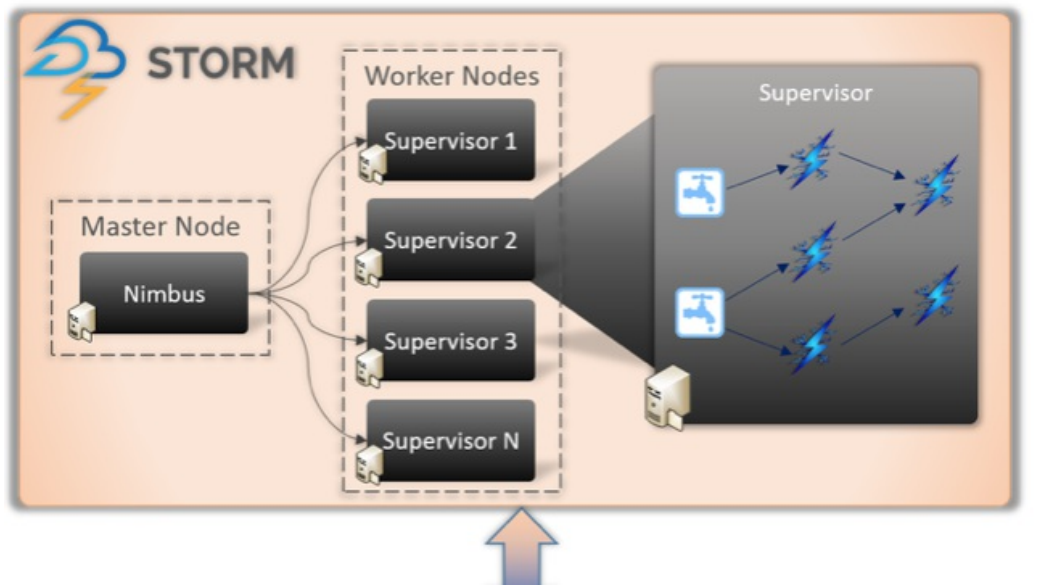
- 시간을 축으로 하여 계속해서 생성되는 데이터
- 배치 분석 vs. 실시간 분석
- 마케팅과 같이 실시간 대응이 중요한 경우에는 데이터 스트림을 실시간 으로 분석
- EX> 이벤트 감지를 통한 이상 거래 탐지(Fraud Detection)

### Storm 특징

- Integrates
  - Storm integrates with any queueing system and any database system.
- Simple API
- Scalable
  - Storm was benchmarked at processing one million 100 byte messages per second per node on hardware with the following specs:
    - Processor: 2x Intel E5645@2.4Ghz
    - Memory: 24 GB
- Fault tolerant
- Guarantees data processing
  - Using Trident, a higher level abstraction over Storm's basic abstractions, you can achieve exactly-once processing semantics.
- Use with any language
- Easy to deploy and operate
- Free and open source


- Apache Storm is a free and open source project licensed under the Apache License, Version 2.0

## Storm Architecture



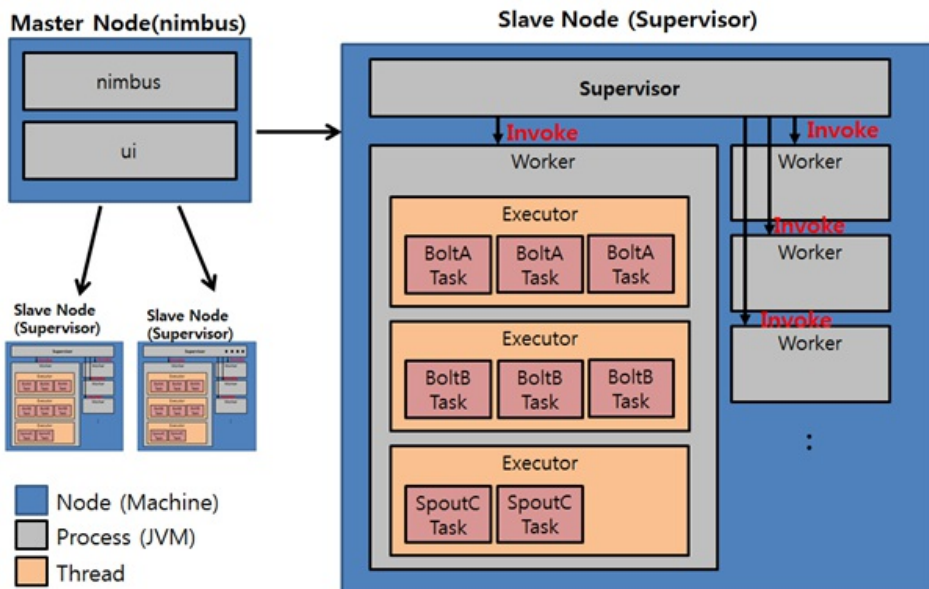
1. Nimbus -> Master Node
2. Supervisor -> Worker Node

## Storm Components



1. Spout: Spout is the entry point in a storm topology.
2. Bolt: Bolt contains the actual processing logic.
3. Topology: Topology is a network of spouts and bolts
4. Stream: A stream is an unbounded sequence of tuples
5. Tuple: Tuple is the most basic data structure in storm . Its a named list of values.

## Storm Parallelism



- Node
  - Nimbus나 Supervisor 프로세스가 기동되는 물리적인 서버
- Worker
  - Supervisor가 있는 노드에서 기동되는 자바 프로세스
- Executor
  - Worker내에서 수행되는 자바 쓰레드
- Task
  - Bolt 및 Spout 객체

## Storm Example

### 1. HelloTopology Example

pom.xml

```
<dependency>
  <groupId>org.apache.storm</groupId>
  <artifactId>storm-core</artifactId>
  <version>1.1.1</version>
  <scope>provided</scope>
</dependency>
```

HelloSpout.java

```
public class HelloSpout extends BaseRichSpout {
    private SpoutOutputCollector collector;

    public void open(Map conf, TopologyContext context,
        SpoutOutputCollector collector) {
        this.collector = collector;
    }
}
```

```

    public void nextTuple() {
        this.collector.emit(new Values("Hello World!"));
    }

    public void declareOutputFields(OutputFieldsDeclarer declarer) {
        declarer.declare(new Fields("say"));
    }
}

```

## HelloBolt.java

```

public class HelloBolt extends BaseBasicBolt {

    private static final long serialVersionUID = 1L;

    public void execute(Tuple tuple, BasicOutputCollector collector) {
        String value = tuple.getStringByField("say");
        System.out.println("Tuple value is " + value);
    }

    public void declareOutputFields(OutputFieldsDeclarer declarer) {
    }
}

```

## HelloTopology.java

```

public class HelloTopology {
    public static void main(String args[]) throws Exception {
        TopologyBuilder builder = new TopologyBuilder();
        builder.setSpout("HelloSpout", new HelloSpout(), 2);
        builder.setBolt("HelloBolt", new HelloBolt(),
4).shuffleGrouping("HelloSpout");

        Config conf = new Config();
        conf.setDebug(true);

        // If there are arguments, we are running on a cluster
        if (args != null && args.length > 0) {
            // parallelism hint to set the number of workers
            conf.setNumWorkers(3);
            // submit the topology
            StormSubmitter.submitTopology(args[0], conf,
builder.createTopology());
        }
        // Otherwise, we are running locally
        else {
            LocalCluster cluster = new LocalCluster();

```

```

        cluster.submitTopology("HelloTopologyLocal", conf,
builder.createTopology());
        Utils.sleep(10000);

        // kill the LearningStormTopology
        cluster.killTopology("HelloTopologyLocal");
        // shutdown the storm test cluster
        cluster.shutdown();
    }
}
}

```

## 2. 빌드 및 배포

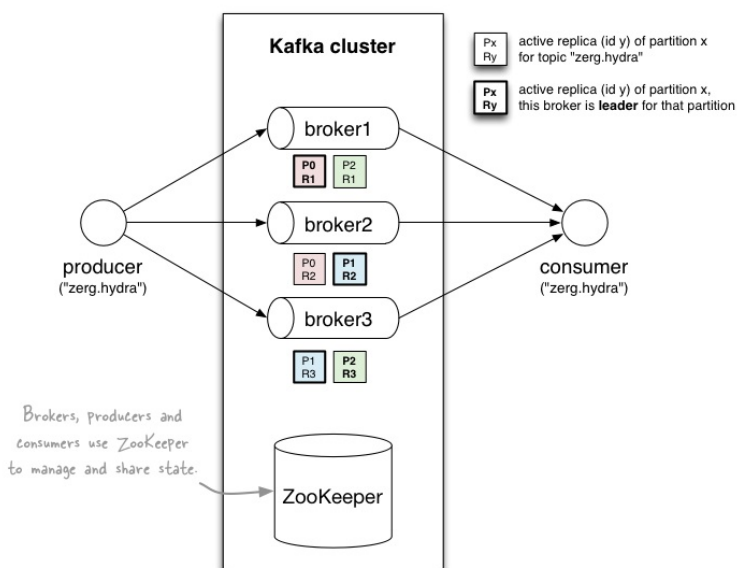
```
$ mvn clean package -DskipTests=true
```

```
$ storm jar storm-0.0.1-SNAPSHOT.jar hello.HelloTopology HelloTopology
```

## Kafka 및 Redis 활용

### 4.1 Kafka

A distributed streaming platform



First a few concepts:


- Kafka is run as a cluster on one or more servers.
- The Kafka cluster stores streams of records in categories called topics.



- Each record consists of a key, a value, and a timestamp.

## Topics and Logs

### Anatomy of a Topic



## Producers

## Consumers

## 4.2 Kafka-Storm Example

pom.xml

```
<dependency>
  <groupId>org.apache.storm</groupId>
  <artifactId>storm-core</artifactId>
  <version>1.1.1</version>
  <scope>provided</scope>
</dependency>
```

CutLogBolt.java

```
public class CutLogBolt extends BaseBasicBolt {

    private static final long serialVersionUID = 1L;

    @Override
    public void execute(Tuple input, BasicOutputCollector collector) {

        String[] splitArray = input.getString(0).split(";");
        String key = "";
        String doctype = "";

        for (int i = 0; i < splitArray.length; i++) {
```

```

        if (splitArray[i].contains("key"))
            key = splitArray[i];
        if (splitArray[i].contains("doctype"))
            doctype = splitArray[i];
    }
    collector.emit(new Values(key, doctype));
}

@Override
public void declareOutputFields(OutputFieldsDeclarer declarer) {
    declarer.declare(new Fields("key", "doctype"));
}
}

```

### ClassifyKeyBolt.java

```

public class ClassifyKeyBolt extends BaseBasicBolt {

    private static final long serialVersionUID = 1L;

    @Override
    public void execute(Tuple input, BasicOutputCollector collector) {
        String[] splitdoctype =
input.getStringByField("doctype").split(":");
        String[] splitkey = input.getStringByField("key").split(":");
        if (splitkey.length == 2 && splitdoctype.length == 2) {
            String doctype = splitdoctype[1].trim();
            String key = splitkey[1].trim();
            // System.err.println(key + ":" + doctype);
            collector.emit(new Values(key + ":" + doctype));
        }
    }

    @Override
    public void declareOutputFields(OutputFieldsDeclarer declarer) {
        declarer.declare(new Fields("subdoctype"));
    }
}

```

### DoctypeCountBolt.java

```

public class DoctypeCountBolt extends BaseBasicBolt {

    private static final long serialVersionUID = 1L;

    Map<String, Integer> docMap = new HashMap<String, Integer>();

    @Override

```

```

    public void execute(Tuple input, BasicOutputCollector collector) {

        String doctype = input.getStringByField("subdoctype");

        Integer count = docMap.get(doctype);
        if (count == null)
            count = 0;

        count++;

        docMap.put(doctype, count);
        System.out.println(docMap);
        collector.emit(new Values(docMap));
    }

    @Override
    public void declareOutputFields(OutputFieldsDeclarer declarer) {
        declarer.declare(new Fields("docmap"));
    }
}

```

### MonitorTopology.java

```

public class MonitorTopology {

    public static void main(String[] args) throws Exception {

        String zkUrl = "demo:2181";
        String brokerUrl = "demo:9092";

        System.out.println("Using Kafka zookeeper url: " + zkUrl + "
broker url: " + brokerUrl);

        ZkHosts hosts = new ZkHosts(zkUrl);

        SpoutConfig spoutConfig = new SpoutConfig(hosts, "onlytest",
"/onlytest", UUID.randomUUID().toString());
        spoutConfig.scheme = new SchemeAsMultiScheme(new
StringScheme());
        KafkaSpout kafkaSpout = new KafkaSpout(spoutConfig);

        TopologyBuilder builder = new TopologyBuilder();
        builder.setSpout("spout", kafkaSpout, 1);
        builder.setBolt("cutbolt", new CutLogBolt(),
2).shuffleGrouping("spout");
        builder.setBolt("classifybolt", new ClassifyKeyBolt(),
2).fieldsGrouping("cutbolt", new Fields("key", "doctype"));
        builder.setBolt("docbolt", new DoctypeCountBolt(),
2).fieldsGrouping("classifybolt", new Fields("subdoctype"));
    }
}

```

```

        Config conf = new Config();
        conf.setDebug(true);

        List<String> nimbus_seeds = new ArrayList<String>();
        nimbus_seeds.add("nimbus url");
        if (args != null && args.length > 0) {
            conf.setNumWorkers(3);
            StormSubmitter.submitTopologyWithProgressBar(args[0],
conf, builder.createTopology());
        } else {
            conf.setNumWorkers(3);
            StormSubmitter.submitTopology("onlytest", conf,
builder.createTopology());
        }
    }
}

```

빌드 및 배포

```

$ mvn clean package -DskipTests=true
$ storm jar storm-0.0.1-SNAPSHOT.jar monitor.MonitorTopology MonitorTopology

```

## 4.2 Redis(Remote Dictionary Server)

Redis is an open source (BSD licensed), **in-memory data structure store**, used as a database, **cache** and message broker

특징

- key-value data store
- value can be a string, list, set, sorted set or hash

Redis Java Example

pom.xml

```

<dependency>
  <groupId>redis.clients</groupId>
  <artifactId>jedis</artifactId>
  <version>2.9.0</version>
</dependency>

```

HelloRedis.java

```

public class HelloRedis {

```

```

    public static void main(String args[]) {
        JedisPool jedisPool = new JedisPool(new JedisPoolConfig(),
"192.168.33.10");
        Jedis jedis = jedisPool.getResource();

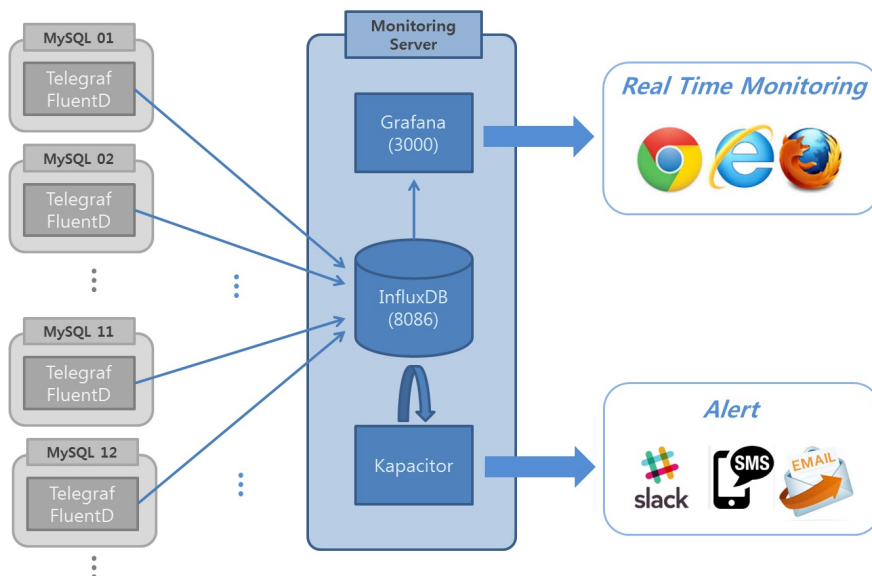
        try {
            jedis.set("test1", "testValue1");
            jedis.set("test2", "testValue2");
            System.out.println("jedis Value 1 : " +
jedis.get("test1"));
            System.out.println("jedis Value 2 : " +
jedis.get("test2"));
            System.out.println(jedis.dbSize());
        } catch (JedisConnectionException e) {
        } finally {
            if (null != jedis) {
                jedisPool.close();
            }
        }

        jedisPool.destroy();
    }
}

```

## Grafana 와 InfluxDB를 이용한 대시보드 구성

### Dashboard Architecture



[출처링크] [http://www.popit.kr/influxdb\\_telegraf\\_grafana\\_2/](http://www.popit.kr/influxdb_telegraf_grafana_2/)

## Grafana

The open platform for beautiful analytics and monitoring

The leading open source software for the time series analytics

## InfluxDB

InfluxDB is a Time Series Database built from the ground up to handle high write & query loads.

## Telegraf

Telegraf is the Agent for Collecting & Reporting Metrics & Data

## Reference

- [Vagrant Getting Started](#)
- [MapReduce Tutorial](#)
- [Apache HBase Reference Guild](#)
- [Apache Storm Documentation, Version: 1.1.1](#)
- [Apache Storm을 이용한 실시간 데이터 처리](#)
- [Spring Data Redis](#)
- <http://lambda-architecture.net/>
- [Zookeeper Overview](#)
- [HBase with Java](#)
- [Apache Kafka Introduction](#)